

Spark SQL


and a bit of partitioning

Anatolii Bardukov
YDATA 2022

Plan

1. RDD Joins
2. Partitions over RDD
3. SQL basics
4. SQL-on-Hadoop
5. Spark SQL
 - a. simple selection
 - b. functions
 - c. joins
 - d. aggregation
 - e. UDF
 - f. Window Function

Book to read




Hadoop Partitions

Block Replication

Namenode (Filename, numReplicas, block-ids, ...)
/users/sameerp/data/part-0, r:2, {1,3}, ...
/users/sameerp/data/part-1, r:3, {2,4,5}, ...

Datanodes



RDD Partitions

Same idea


RDD is split into partitions

Possible to split RDD into partitions by key

Different partitions are consumed by different Spark workers

Dataset is broken into partitions

Partitions are each stored in a worker's memory



RDD Partitions management

We can not tell exactly, where each partition goes in terms of nodes - inefficient

But we can create partition function and amount of buckets

Partition tuning example

```
numbers = sc.parallelize(range(10))
```

```
squares = numbers.map(lambda x: (x, x**2))
```

```
def custom_partitioner(value):  
    return value % 3
```

```
squares.partitionBy(3, custom_partitioner).glom().collect()
```

```
[[ (0, 0), (6, 36), (9, 81), (3, 9) ],  
 [ (1, 1), (4, 16), (7, 49) ],  
 [ (5, 25), (2, 4), (8, 64) ]]
```

RDD Joins

Transformation to join several different RDDs into single one using some key such that key remains same and all other rows are combined into the value

```
>>> x = sc.parallelize([("a", 1), ("b", 4)])  
>>> y = sc.parallelize([("a", 2), ("a", 3)])  
>>> sorted(x.join(y).collect())  
[('a', (1, 2)), ('a', (1, 3))]
```


Advanced Partition Tuning example




Figure 4-4. Each join of `userData` and `events` without using `partitionBy()`

Advanced Partition Tuning example


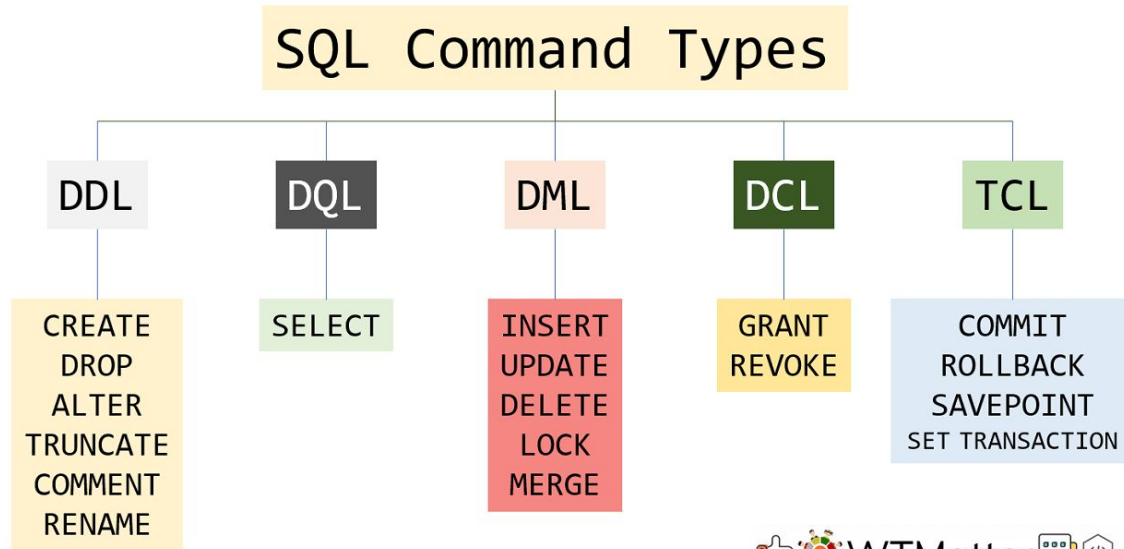


Figure 4-5. Each join of `userData` and `events` using `partitionBy()`

SQL

SQL

Structured Query Language - **declarative** domain-specific language usually used in (relational) DBMS, takes advantage of **relations** within the **structured** data



SQL ANSI standard

Consists of 10 parts, latest release SQL:2016, SQL 2019 (adds part 15, MDarrays)

<https://www.iso.org/standard/63555.html> - first part

Usually, DBMS implement SQL features before they go to standard:

Oracle: PL/SQL

Postgres: PL/pgSQL

SQL syntax difference example

https://www.w3schools.com/sql/sql_top.asp

SELECT **TOP** n from ... - *MS Access*

SELECT * from ... **LIMIT** N - *MySQL*

SELECT * from ... WHERE **ROWNUM** <= n - *Oracle*

DDL

Data **Definition** Language

Subset of SQL responsible for definition of data structures

Commands: CREATE/ALTER/DROP database, table, index, view

Exists in Spark SQL

DDL

```
CREATE TABLE [table name] ( [column definitions] ) [table parameters]
```

```
CREATE TABLE employees (
```

```
    id            INTEGER      PRIMARY KEY,
```

```
    first_name    VARCHAR(50)  not null,
```

```
    last_name     VARCHAR(75)  not null,
```

```
    fname        VARCHAR(50)   not null,
```

```
    dateofbirth   DATE         not null
```

```
);
```


DQL

Data **Query** Language

(possible alternatives: Data Retrieval Language/Data Selection Language)

Describes details of SELECT statement syntax, used to retrieve data from database

Usually treated as a part of DML

Exists in Spark SQL

DQL

```
SELECT [ hints , ... ] [ ALL | DISTINCT ] { named_expression [ , ... ] }  
    FROM { from_item [ , ... ] }  
    [ WHERE boolean_expression ]  
    [ GROUP BY expression [ , ... ] ]  
    [ HAVING boolean_expression ]
```

```
[ WITH with_query [ , ... ] ]  
select statement [ { UNION | INTERSECT | EXCEPT } [ ALL | DISTINCT ] select statement, ... ]  
    [ ORDER BY { expression [ ASC | DESC ] [ NULLS { FIRST | LAST } ] [ , ... ] } ]  
    [ SORT BY { expression [ ASC | DESC ] [ NULLS { FIRST | LAST } ] [ , ... ] } ]  
    [ CLUSTER BY { expression [ , ... ] } ]  
    [ DISTRIBUTE BY { expression [ , ... ] } ]  
    [ WINDOW { named_window [ , WINDOW named_window, ... ] } ]  
    [ LIMIT { ALL | expression } ]
```

```
SELECT name, age FROM person ORDER BY age DESC NULLS FIRST;
```

DML

Data **Manipulation** Language

Describes set of commands responsible for updates of the data inside database,
i.e.:

INSERT / UPDATE / DELETE

Partially exists in Spark SQL

DML

```
INSERT INTO [ TABLE ] table_identifier [ partition_spec ]  
    { VALUES ( { value | NULL } [ , ... ] ) [ , ( ... ) ] | query }
```

```
INSERT INTO students VALUES ('Amy Smith', '123 Park Ave, San Jose', 111111);
```

DCL

Data **Control** Language

Used to control access to data and manipulate permissions:

GRANT / DENY / REVOKE

Does **not** exists in Spark SQL

DCL

Postgres Example

Revoke insert privilege for the public on table `films`:

```
REVOKE INSERT ON films FROM PUBLIC;
```

Revoke all privileges from user `manuel` on view `kinds`:

```
REVOKE ALL PRIVILEGES ON kinds FROM manuel;
```

TCL

Transaction Control Language

Used to manage transactions in the DB:

COMMIT / ROLLBACK / SAVEPOINT

SQL on Hadoop

Shark (SQL on Spark)

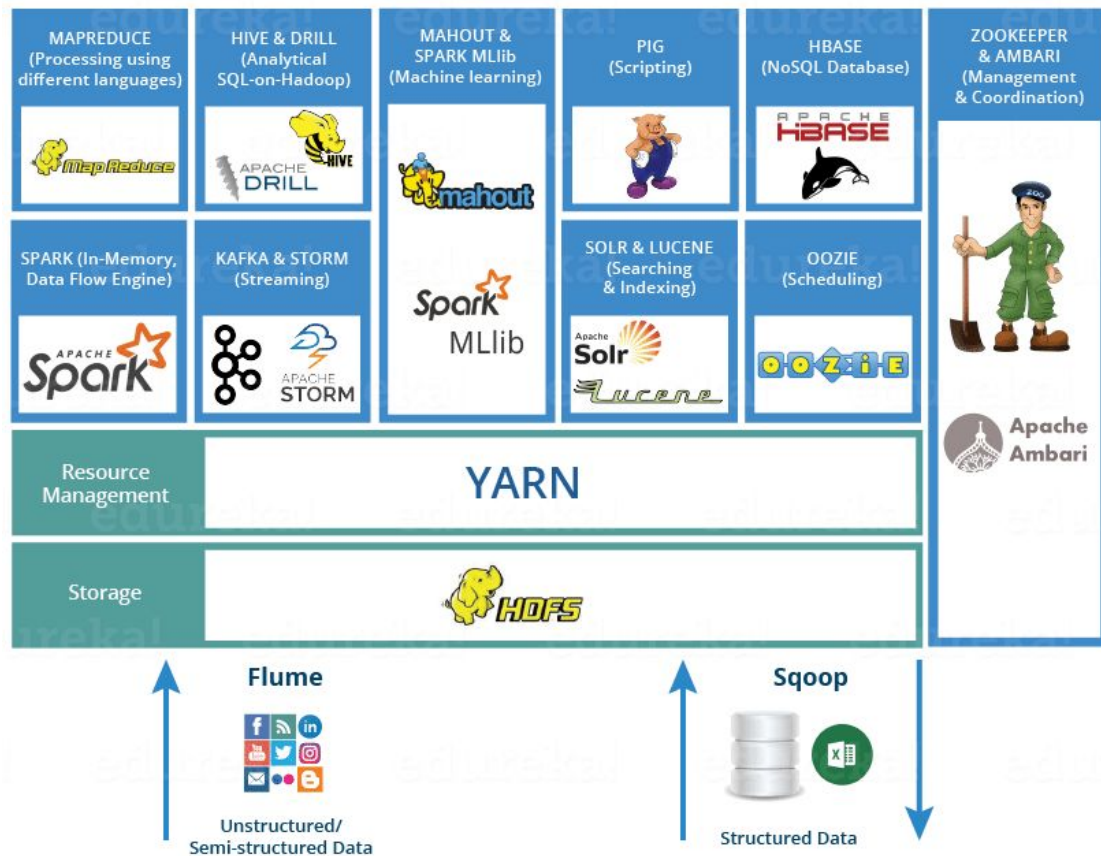
Spark SQL

Hive

Impala

Presto

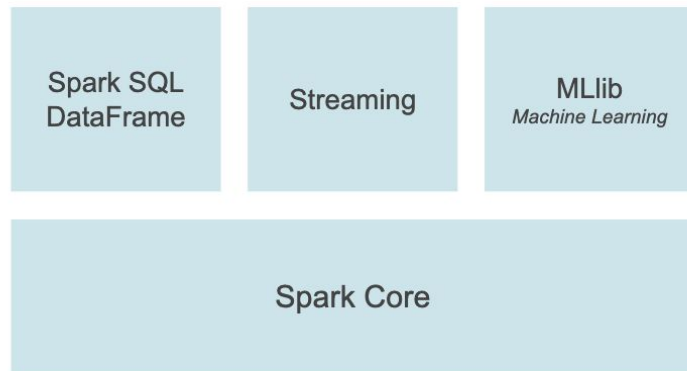
...



Spark SQL

Spark SQL is a Spark module for structured data processing

It provides a programming abstraction called DataFrames and can also act as a distributed SQL query engine



Spark SQL architecture

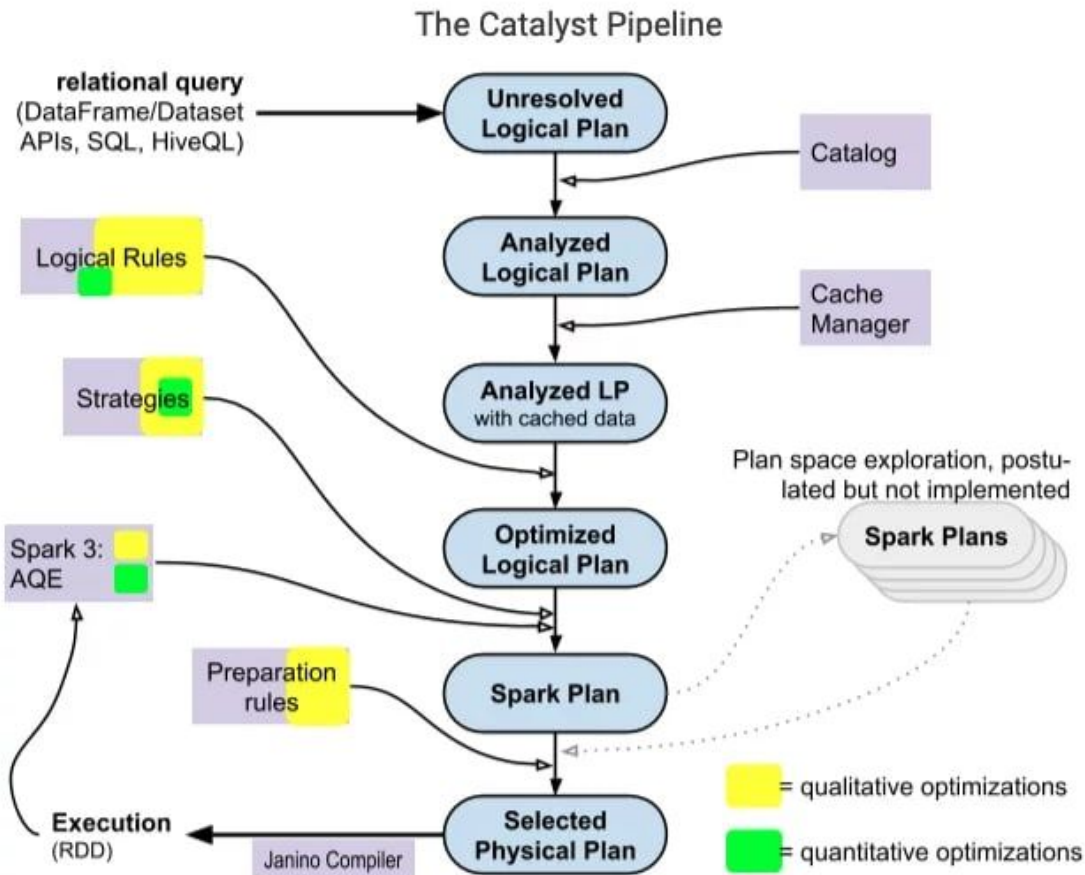
DataFrame API / SQL query goes to planner

Optimizations performed

Logical plan becomes Physical

Plan is materialized to RDD

Adaptive Query Execution optimizes the query while run



Spark SQL types

Before we continue with Spark SQL syntax and commands, data types to work with:

<https://spark.apache.org/docs/latest/sql-ref-datatypes.html>

- Numeric types
- String type
- Binary type
- Boolean type
- Datetime type
- Complex types

Spark SQL types

```
SELECT double('infinity') * 0 AS col;
```

```
+----+
```

```
|col|
```

```
+----+
```

```
|NaN|
```

```
+----+
```

```
SELECT double('-infinity') * (-1234567) AS col;
```

```
+-----+
```

```
|      col|
```

```
+-----+
```

```
|Infinity|
```

```
+-----+
```

Spark SQL syntax

<https://spark.apache.org/docs/latest/sql-ref-syntax.html>

- Data Definition Statements
- Data Manipulation Statements
- Data Retrieval (Queries)
- Auxiliary Statements

Spark SQL SELECT

Probably, most useful statement, used to query data from tables:

```
SELECT 10;
```

```
SELECT 10 as column_name;
```

```
SELECT "value";
```

```
SELECT col_name from table_name;
```

Has a lot of additional params

Spark SQL WHERE

Filtering condition (boolean expression) for data coming from FROM section of SELECT statement

```
SELECT  
  col_name  
FROM  
  table_name  
WHERE  
  col_name > 10
```

SQL small task (table name: books)

ID	author_id	title	language
1	1	Eugene Onegin	RU
2	2	The Lord of the Rings I	ENG
3	2	The Lord of the Rings II	ENG

SELECT

book_id	title
2	The Lord of the Rings I
3	The Lord of the Rings II

Answer

```
SELECT id AS "book_id", title  
FROM books  
WHERE title like '%Lord%'
```

Spark SQL Functions

We can use different functions in select statement

Spark defines a set of built-in function and supports user-defined functions (UDFs)

Those functions would be applied to data rows

Different types:

- simple, i.e. *lower* for string data
- specific for JSON/Datetime/Map/Array, i.e. *map_concat*
- aggregation/window, i.e. *min*, *max*, *mean* for numeric values

Spark SQL Functions

Analytical functions are applied to the group of rows:

```
SELECT count(*) FROM table_name
```

would return count of rows from table

Spark SQL GROUP BY

Aggregation function (like *count* we had previously) can be used for group of rows.

GROUP BY clause is used to create such groups, where aggregation functions would be applied

For example, we can calculate amount of books written by the author:

```
SELECT author, count(*) FROM books GROUP BY author
```

Spark SQL HAVING

Used to filter the results produced by GROUP BY based on the specified condition

Just like WHERE condition, but after GROUP BY is applied

How can we select only authors, who have written more than 100 books?

```
SELECT author, count(*) as books_cnt  
FROM books
```

...

Spark SQL HAVING

Used to filter the results produced by GROUP BY based on the specified condition

Just like WHERE condition, but after GROUP BY is applied

How can we select only authors, who have written more than 100 books:

```
SELECT author, count(*) as books_cnt
FROM books
GROUP BY author
HAVING books_cnt > 100
```

Spark SQL naming and aliasing

In previous example we used naming of the column:

```
SELECT 10 as col_name;
```

We can also use aliases for tables:

```
SELECT t.* from table t;
```

Spark SQL ORDER BY

Used to return the result rows in a sorted manner in the specified order

Guarantees a total order in the output

For example, we can order books by the year there were written:

```
SELECT name, year FROM books ORDER BY year DESC
```


Spark SQL SORT BY

Used to return the result rows sorted within each partition in the specified order

When there is more than one partition SORT BY may return result that is partially ordered

Different than ORDER BY clause which guarantees a total order of the output

Spark SQL EXPLAIN

Statement is used to provide logical/physical plans for an input statement

By default, this clause provides information about a physical plan only, but can be
EXTENDED

Spark SQL EXPLAIN

-- Default Output


```
EXPLAIN select k, sum(v) from values (1, 2), (1, 3) t(k, v) group by k;
```

```
+-----+
|                                             plan|
+-----+
| == Physical Plan ==
*(2) HashAggregate(keys=[k#33], functions=[sum(cast(v#34 as bigint))])
+- Exchange hashpartitioning(k#33, 200), true, [id=#59]
   +- *(1) HashAggregate(keys=[k#33], functions=[partial_sum(cast(v#34 as bigint))])
      +- *(1) LocalTableScan [k#33, v#34]
|
+-----+
```

-- Using Extended

```
EXPLAIN EXTENDED select k, sum(v) from values (1, 2), (1, 3) t(k, v) group by k;
```

Spark SQL JOIN



Spark SQL JOIN

ID	Name
1	Pushkin
2	Tolkien

Author ID	ID	Title	Language
1	1	Eugene Onegin	RU
2	2	The Lord of the Rings I	ENG
2	3	The Lord of the Rings II	ENG

Join on Author ID

Author ID	ID	Name	Title	Language
1	1	Pushkin	Eugene Onegin	RU
2	2	Tolkien	The Lord of the Rings I	ENG
2	3	Tolkien	The Lord of the Rings II	ENG

Spark SQL Join

-- Use employee and department tables to demonstrate left join.

```
SELECT id, name, employee.deptno, deptname  
      FROM employee LEFT JOIN department ON employee.deptno = department.deptno;
```

id	name	deptno	deptname
105	Chloe	5	NULL
103	Paul	3	Engineering
101	John	1	Marketing
102	Lisa	2	Sales
104	Evan	4	NULL
106	Amy	6	NULL

Harder question

id	name
1	Bob
2	Marta

author_id	id	title	language
1	1	Eugene Onegin	RU
2	2	The Lord of the Rings I	ENG
2	3	The Lord of the Rings II	ENG

id	member_id	book_id	date
1	1	1	01.02.2020
2	1	2	02.03.2020
3	2	1	07.11.2020
4	2	2	21.07.2019
5	2	3	27.08.2019

JOIN

name	title	language	...
Bob	Eugene Onegin	RU	...
Bob	The Lord of the Rings I	ENG	...
Marta	Eugene Onegin	RU	...
Marta	The Lord of the Rings I	ENG	...
Marta	The Lord of the Rings II	ENG	...

GROUP

Full name	# of books borrowed
Bob	1
Marta	2



Answer

```
SELECT
    m.name AS "Full Name",
    count(*) AS "# of books borrowed"
FROM borrowings br
JOIN books b ON br.book_id = b.id
JOIN members m ON br.member_id = m.id
WHERE b.language = 'EN'
GROUP BY m.name;
```


Spark SQL Window Function

Operate on a group of rows, referred to as a window, and calculate a return value for each row based on the group of rows

Window functions are useful for processing tasks such as calculating a moving average, computing a cumulative statistic, or accessing the value of rows given the relative position of the current row.



Spark SQL Window Function

```
SELECT name, salary,  
       LAG(salary) OVER (PARTITION BY dept ORDER BY salary) AS lag,  
       LEAD(salary, 1, 0) OVER (PARTITION BY dept ORDER BY salary) AS lead  
FROM employees;
```

name	dept	salary	lag	lead
Lisa	Sales	10000	NULL	30000
Alex	Sales	30000	10000	32000
Evan	Sales	32000	30000	0
Fred	Engineering	21000	NULL	23000
Chloe	Engineering	23000	21000	23000
Tom	Engineering	23000	23000	29000
Paul	Engineering	29000	23000	0
Helen	Marketing	29000	NULL	29000
Jane	Marketing	29000	29000	35000
Jeff	Marketing	35000	29000	0

Spark SQL Window Function

Examples of functions, that can be used within window frame:

Ranking Functions

Syntax: RANK | DENSE_RANK | PERCENT_RANK | NTILE | ROW_NUMBER

Analytic Functions

Syntax: CUME_DIST | LAG | LEAD

Aggregate Functions

Syntax: MAX | MIN | COUNT | SUM | AVG | ...

Window Function question

How to get top N (by salary) employees from each department?

table: Employee(name: string, salary: int, department: string)

Window Function question

How to get top N (by salary) employees from each department?

table: Employee(name: string, salary: int, department: string)

```
SELECT * FROM (  
    SELECT e.*,  
    ROW_NUMBER() OVER (PARTITION BY department ORDER BY salary DESC) rn  
    FROM Employee e  
)  
WHERE rn <= N
```